

Data Preprocessing

Import the necessary libraries

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
```

Read the data

```
bids_df = pd.read_csv('bids.csv')
train_df = pd.read_csv('train.csv')
test_df = pd.read_csv('test.csv') # Count the number of bids from each country
```

```
country_counts = bids_df['country'].value_counts()
```

Create a bar plot to visualize the distribution of bids by country


```
fig, ax = plt.subplots(figsize=(18, 4.8))
countries_plot = sns.barplot(x=country_counts.index,
                             y=country_counts.values, ax=ax)
```

Hide the x-axis labels to improve readability

```
ax.get_xaxis().set_visible(False)
```

Display the plot

```
plt.show()
```



```
bids_df.head()
```

	bid_id	bidder_id	auction	merchandise
device \				
0	0	8dac2b259fd1c6d1120e519fb1ac14fbqvax8	ewmzr	jewelry
phone0				
1	1	668d393e858e8126275433046bbd35c6tywop	aeqok	furniture
phone1				
2	2	aa5f360084278b35d746fa6af3a7a1a5ra3xe	wa00e	home goods
phone2				
3	3	3939ac3ef7d472a59a9c5f893dd3e39fh9ofi	jefix	jewelry
phone4				
4	4	8393c48eaf4b8fa96886edc7cf27b372dsibi	jefix	jewelry
phone5				

	time	country	ip	url
0	9759243157894736	us	69.166.231.58	vasstdc27m7nks3
1	9759243157894736	in	50.201.125.84	jmqhflrzway9c
2	9759243157894736	py	112.54.208.157	vasstdc27m7nks3
3	9759243157894736	in	18.99.175.133	vasstdc27m7nks3
4	9759243157894736	in	145.138.5.37	vasstdc27m7nks3

```
train_df.head()
```

	bidder_id	\
0	91a3c57b13234af24875c56fb7e2b2f4rb56a	
1	624f258b49e77713fc34034560f93fb3hu3jo	
2	1c5f4fc669099bfbfac515cd26997bd12ruaj	
3	4bee9aba2abda51bf43d639013d6efe12iycd	
4	4ab12bc61c82ddd9c2d65e60555808acqgos1	

	payment_account	\
0	a3d2de7675556553a5f08e4c88d2c228754av	
1	a3d2de7675556553a5f08e4c88d2c228v1sga	
2	a3d2de7675556553a5f08e4c88d2c2280cybl	
3	51d80e233f7b6a7dfdee484a3c120f3b2ita8	
4	a3d2de7675556553a5f08e4c88d2c22857ddh	

	address	outcome
0	a3d2de7675556553a5f08e4c88d2c228vt0u4	0.0
1	ae87054e5a97a8f840a3991d12611fdcrfbq3	0.0
2	92520288b50f03907041887884ba49c0cl0pd	0.0
3	4cb9717c8ad7e88a9a284989dd79b98dbevyi	0.0
4	2a96c3ce94b3be921e0296097b88b56a7x1ji	0.0

```
bids_df.shape
```

```
(7656334, 9)
```

```
bids_df.isnull().sum()
```

```
bid_id          0
bidder_id       0
auction         0
merchandise     0
device          0
time            0
country        8859
ip              0
url             0
dtype: int64
```

```
missing_percent = bids_df['country'].isnull().mean()
print(f"Percentage of missing data in country column:
{missing_percent*100: .2f}%")
```

```
Percentage of missing data in country column: 0.12%
```

```

# Import additional libraries
import matplotlib.pyplot as plt
import seaborn as sns
import numpy as np

# Count the number of bids from each country
country_counts = bids_df['country'].value_counts()

# Create a bar plot to visualize the distribution of bids by country
fig, ax = plt.subplots(figsize=(18, 4.8))

# Create a custom color map with a gradient
color_map = plt.get_cmap("Greens")
norm = plt.Normalize(country_counts.values.min(),
country_counts.values.max())

# Generate the bar plot using the custom gradient color
countries_plot = sns.barplot(x=country_counts.index,
y=country_counts.values, ax=ax,

palette=color_map(norm(country_counts.values)))

# Hide the x-axis labels to improve readability
ax.get_xaxis().set_visible(False)

# Add a title to the plot
plt.title("Distribution of Bids by Country")

# Display the plot
plt.show()

```



```

bids_df['country'] =
bids_df['country'].fillna(bids_df['country'].mode()[0])

```

Feature Engineering

```

# Sorting the bids DataFrame based on bidder_id and time
sorted_bids = bids_df.sort_values(['bidder_id', 'time'])
sorted_bids.head(15)

```

	bid_id	bidder_id	auction
merchandise \	7179832	001068c415025a009fee375a12cff4fcnht8y	4ifac
jewelry	1281292	002d229ffb247009810828f648afc2ef593rb	2tdw2
mobile	1281311	002d229ffb247009810828f648afc2ef593rb	2tdw2
mobile	6805028	0030a2dd87ad2733e0873062e4f83954mkj86	obbny
mobile	3967330	003180b29c6a5f8f1d84a6b7b6f7be57tjj1o	obbny
mobile	6166636	003180b29c6a5f8f1d84a6b7b6f7be57tjj1o	cqsh6
mobile	7140567	003180b29c6a5f8f1d84a6b7b6f7be57tjj1o	efh5o
goods	2597846	00486a11dff552c4bd7696265724ff81yeo9v	no958 home
goods	2599005	00486a11dff552c4bd7696265724ff81yeo9v	6plix home
goods	2750709	00486a11dff552c4bd7696265724ff81yeo9v	gst86 home
goods	3062627	00486a11dff552c4bd7696265724ff81yeo9v	9ul86 home
goods	3302219	00486a11dff552c4bd7696265724ff81yeo9v	6plix home
goods	3682352	00486a11dff552c4bd7696265724ff81yeo9v	no958 home
goods	3682395	00486a11dff552c4bd7696265724ff81yeo9v	lx0hm home
goods	3786154	00486a11dff552c4bd7696265724ff81yeo9v	gst86 home

url	device	time	country	ip
7179832	phone561	9706345052631578	bn	139.226.147.115
vasstdc27m7nks3	1281292	9766744105263157	sg	37.40.254.131
vasstdc27m7nks3	1281311	9766744210526315	sg	37.40.254.131
vasstdc27m7nks3	6805028	9704553947368421	ir	21.67.17.162
vnw40k8zzokijsv	3967330	9640018631578947	id	44.241.8.179
sj4jidex850loas	6166636	9700605052631578	id	190.88.89.83
sj4jidex850loas	7140567	9705974315789473	id	115.47.140.180
vasstdc27m7nks3	2597846	9632636526315789	ng	143.118.40.162

```

vasstdc27m7nks3
2599005    phone4  9632641157894736      ng    143.118.40.162
vasstdc27m7nks3
2750709    phone45 9633339684210526      ng    54.212.177.220
0wfuw1acucr1cdl
3062627    phone45 9635439947368421      ng    236.63.15.129
n01mdz1usu12kso
3302219    phone788 9636501894736842      ng    127.247.172.237
n01mdz1usu12kso
3682352    phone45 9637993315789473      ng    236.63.15.129
vasstdc27m7nks3
3682395    phone45 9637993684210526      ng    236.63.15.129
vasstdc27m7nks3
3786154    phone4  9638911421052631      ng    220.193.41.160
mm1kg05ew2b5wk6

```

Calculating the time differences between consecutive bids for every bidder

```

time_differences = sorted_bids.groupby('bidder_id')[['time']].diff()
time_differences.head(15)

```

```

              time
7179832         NaN
1281292         NaN
1281311  1.052632e+08
6805028         NaN
3967330         NaN
6166636  6.058642e+13
7140567  5.369263e+12
2597846         NaN
2599005  4.631579e+09
2750709  6.985263e+11
3062627  2.100263e+12
3302219  1.061947e+12
3682352  1.491421e+12
3682395  3.684211e+08
3786154  9.177368e+11

```

Adding the calculated time differences to the sorted_bids DataFrame

```

sorted_bids['time_difference'] = time_differences
time_diff_features = sorted_bids[['bidder_id',
'time_difference']].dropna()
time_diff_features.head()

```

```

              bidder_id  time_difference
1281311  002d229ffb247009810828f648afc2ef593rb  1.052632e+08
6166636  003180b29c6a5f8f1d84a6b7b6f7be57tjjlo  6.058642e+13
7140567  003180b29c6a5f8f1d84a6b7b6f7be57tjjlo  5.369263e+12
2599005  00486a11dff552c4bd7696265724ff81yeo9v  4.631579e+09
2750709  00486a11dff552c4bd7696265724ff81yeo9v  6.985263e+11

```

```
# Grouping the bid intervals by bidder_id and computing descriptive statistics
```

```
bid_stats = time_diff_features.groupby('bidder_id')
[['time_difference']].describe().reset_index()
bid_stats.columns = bid_stats.columns.droplevel(level=0)
bid_stats = bid_stats.rename(columns={'': 'bidder_id', 'mean':
'mean_time_diff', 'std': 'std_time_diff', '50%': 'median_time_diff',
'min': 'min_time_diff', 'max': 'max_time_diff'}).fillna(0)
bid_stats['iqr_time_diff'] = bid_stats['75%'] - bid_stats['25%']
bid_stats = bid_stats.drop(['25%', '75%', 'count'], axis=1)
bid_stats.head()
```

	bidder_id	mean_time_diff
std_time_diff \		
0	002d229ffb247009810828f648afc2ef593rb	1.052632e+08
0.000000e+00		
1	003180b29c6a5f8f1d84a6b7b6f7be57tjjlo	3.297784e+13
3.904443e+13		
2	00486a11dff552c4bd7696265724ff81yeo9v	4.018413e+12
1.153730e+13		
3	0051aef3fdeacdadb664b9b3b07e04e4coc6	1.635106e+11
5.770740e+11		
4	0053b78cde37c4384a20d2da9aa4272aym4pb	7.065316e+09
4.784394e+11		

	min_time_diff	median_time_diff	max_time_diff	iqr_time_diff
0	1.052632e+08	1.052632e+08	1.052632e+08	0.000000e+00
1	5.369263e+12	3.297784e+13	6.058642e+13	2.760858e+13
2	3.684211e+08	9.177368e+11	5.094174e+13	1.696763e+12
3	5.263158e+07	2.736842e+09	3.792368e+12	1.228947e+10
4	0.000000e+00	3.684211e+08	5.002753e+13	1.000000e+09

```
import pandas as pd
```

```
# Load the bids dataset
```

```
bids_df = pd.read_csv('bids.csv')
```

```
# Compute bid statistics by bidder
```

```
bid_stats = bids_df.groupby('bidder_id')['bid_id'].agg(['count',
'mean', 'median'])
```

```
# Rename the 'mean' and 'median' columns in the bid_stats DataFrame to avoid conflicts with existing column names in train_set
```

```
bid_stats = bid_stats.rename(columns={'mean': 'bid_mean', 'median':
'bid_median'})
```

```
# Merge the bid_stats DataFrame with train_set and test_set DataFrames
```

```
train_set = pd.merge(train_df, bid_stats, on='bidder_id', how='left')
train_set = train_set.fillna(train_df.median())
```

```
test_set = pd.merge(test_df, bid_stats, on='bidder_id', how='left')
test_set = test_set.fillna(test_df.median())
```

C:\Users\egodd\AppData\Local\Temp\ipykernel_30568\2778142480.py:14:
FutureWarning: The default value of numeric_only in DataFrame.median
is deprecated. In a future version, it will default to False. In
addition, specifying 'numeric_only=None' is deprecated. Select only
valid columns or specify the value of numeric_only to silence this
warning.

```
train_set = train_set.fillna(train_df.median())
```

C:\Users\egodd\AppData\Local\Temp\ipykernel_30568\2778142480.py:17:
FutureWarning: The default value of numeric_only in DataFrame.median
is deprecated. In a future version, it will default to False. In
addition, specifying 'numeric_only=None' is deprecated. Select only
valid columns or specify the value of numeric_only to silence this
warning.

```
test_set = test_set.fillna(test_df.median())
```

*# Calculating the number of simultaneous bids (bids with a time
difference of 0)*

```
simultaneous_bids =  
time_diff_features[time_diff_features['time_difference'] ==  
0].groupby('bidder_id').count().reset_index()  
simultaneous_bids =  
simultaneous_bids.rename(columns={'time_difference':  
'num_simultaneous_bids'})  
simultaneous_bids.head()
```

	bidder_id	num_simultaneous_bids
0	0053b78cde37c4384a20d2da9aa4272aym4pb	728
1	00a79ebd15f0b24a0a3b5794457cd8ed7dng1	29
2	00b519ec8ed5e370328451379bb708a306eoj	1
3	00e0f614d9dd32dd27f6080f472d2934emlos	15
4	019cf2d366df756c092c91e26f406acdozha7	1

*# Merging the simultaneous_bids DataFrame with train_set and test_set
DataFrames*

```
train_set = train_set.merge(simultaneous_bids, on='bidder_id',  
how='left').fillna(0)
```

```
test_set = test_set.merge(simultaneous_bids, on='bidder_id',  
how='left').fillna(0)
```

Total number of bids made by each

```
num_bids = bids_df.groupby('bidder_id')  
['bid_id'].count().reset_index().rename(columns={'bid_id':  
'num_bids'})  
num_bids.head()
```

Merging with train_set and test_set

```
train_set = train_set.merge(num_bids, on='bidder_id',
```

```

how='left').fillna(0)
test_set = test_set.merge(num_bids, on='bidder_id',
how='left').fillna(0)

# Count the number of auctions for each bidder
num_auct = bids_df.groupby('bidder_id')
['auction'].nunique().reset_index()
num_auct = num_auct.rename(columns={'auction': 'num_auct'})

# Merge the num_auct DataFrame with train_set and test_set
train_set = train_set.merge(num_auct, on='bidder_id', how='left')
test_set = test_set.merge(num_auct, on='bidder_id',
how='left').fillna(0)

# Count the number of unique device types for each bidder
num_device_type = bids_df.groupby('bidder_id')
['device'].nunique().reset_index()
num_device_type = num_device_type.rename(columns={'device':
'num_device_type'})

# Merge the num_device_type DataFrame with train_set and test_set
train_set = train_set.merge(num_device_type, on='bidder_id',
how='left')
test_set = test_set.merge(num_device_type, on='bidder_id',
how='left').fillna(0)

# Count the number of unique URLs for each bidder
num_url = bids_df.groupby('bidder_id')['url'].nunique().reset_index()
num_url = num_url.rename(columns={'url': 'num_url'})

# Merge the num_url DataFrame with train_set
train_set = train_set.merge(num_url, on='bidder_id', how='left')
test_set = test_set.merge(num_url, on='bidder_id', how='left')

# Total number of unique IPs and countries for each bidder_id
num_ip_ctry = bids_df.groupby('bidder_id')[['ip',
'country']].nunique().reset_index().rename(columns={'ip': 'num_ip',
'country': 'num_ctry'})

# Merging with train_set and test_set
train_set = train_set.merge(num_ip_ctry, on='bidder_id',
how='left').fillna(0)
test_set = test_set.merge(num_ip_ctry, on='bidder_id',
how='left').fillna(0)

# Maximum number of bids per device for each bidder_id
max_bids_per_device = bids_df.groupby(['bidder_id', 'device'])
['bid_id'].count().reset_index().groupby('bidder_id')
['bid_id'].max().reset_index().rename(columns={'bid_id':
'max_bids_per_device'})

```



```

# Merging with train_set and test_set
train_set = train_set.merge(max_bids_per_device, on='bidder_id',
how='left').fillna(0)
test_set = test_set.merge(max_bids_per_device, on='bidder_id',
how='left').fillna(0)

# Sorting bids by auction and time, then counting the number of first
bids for each bidder
sorted_first_bids = bids_df.sort_values(['auction', 'time'])
first_bids =
sorted_first_bids.groupby('auction').first().reset_index()
first_bids_count = first_bids.groupby('bidder_id').count()
['bid_id'].reset_index()
first_bids_count = first_bids_count.rename(columns={'bid_id':
'num_first_bids'})
first_bids_count.head()

```

	bidder_id	num_first_bids
0	0053b78cde37c4384a20d2da9aa4272aym4pb	4
1	00a79ebd15f0b24a0a3b5794457cd8ed7dng1	7
2	00e0f614d9dd32dd27f6080f472d2934emlos	6
3	019cf2d366df756c092c91e26f406acdozha7	2
4	01cda526658455000913950f20cf31a2q6nsf	8

```

# Sorting bids by auction and time in descending order, then counting
the number of last bids for each bidder
sorted_last_bids = bids_df.sort_values(['auction', 'time'],
ascending=[True, False])
last_bids = sorted_last_bids.groupby('auction').first().reset_index()
last_bids_count = last_bids.groupby('bidder_id').count()
['bid_id'].reset_index()
last_bids_count = last_bids_count.rename(columns={'bid_id':
'num_last_bids'})
last_bids_count.head()

```

	bidder_id	num_last_bids
0	0053b78cde37c4384a20d2da9aa4272aym4pb	6
1	00b519ec8ed5e370328451379bb708a306eoj	1
2	00e0f614d9dd32dd27f6080f472d2934emlos	3
3	01067975436d123f717ee5aba0dd4bbfa0937	1
4	0113d101ec6aabd354adac645a1ec3e82ln88	1

```

# Check if 'bidder_id' exists in train_set dataframe
if 'bidder_id' in train_set.columns:
    # Merging the first_bids_count and last_bids_count DataFrames with
the train_set and test_set DataFrames
    train_set = train_set.merge(first_bids_count, on='bidder_id',
how='left').fillna(0)
    train_set = train_set.merge(last_bids_count, on='bidder_id',
how='left').fillna(0)

```

```

# Check if 'bidder_id' exists in test_set dataframe
if 'bidder_id' in test_set.columns:
    test_set = test_set.merge(first_bids_count, on='bidder_id',
how='left').fillna(0)
    test_set = test_set.merge(last_bids_count, on='bidder_id',
how='left').fillna(0)

# Calculate auction duration
auct_duration = bids_df.sort_values(['auction', 'time'])[['bidder_id',
'auction', 'time']]
auct_duration =
auct_duration[['auction', 'time']].groupby('auction').agg([max,min]).re
set_index().droplevel(axis=1, level=0).rename(columns={'': 'auction'})
auct_duration['auct_duration'] = auct_duration['max'] -
auct_duration['min']
auct_duration.head()

    auction          max          min  auct_duration
0    00270  9709212894736842  9699049894736842  101630000000000
1    008vv  9760397157894736  9759369421052631  1027736842105
2    00cwr  9698636578947368  9695641631578947  2994947368421
3    00do0  9759865210526315  9759323842105263  541368421052
4    00hgy  9772723842105263  9759410368421052  13313473684211

# Calculate bid ratios in the first and second half of auctions
time_ratio = bids_df.sort_values(['auction', 'time'])[['bidder_id',
'auction', 'time']]
time_ratio = time_ratio.merge(auct_duration, on='auction', how='left')
time_ratio['temp'] = time_ratio['time'] - time_ratio['auct_duration']
/ 2
time_ratio['firsthalf'] = time_ratio['temp'] < time_ratio['min']

# Aggregate bid ratios
ratio_firsthalf = time_ratio[['bidder_id',
'firsthalf']].groupby('bidder_id').agg(['count',
sum]).reset_index().droplevel(axis=1, level=0).rename(columns={'':
'bidder_id', 'count': 'num_total_bids', 'sum': 'num_firsthalf_bids'})
ratio_firsthalf['num_secondhalf_bids'] =
ratio_firsthalf['num_total_bids'] -
ratio_firsthalf['num_firsthalf_bids']

print(train_set.columns)

Index(['bidder_id', 'payment_account', 'address', 'outcome', 'count',
'bid_mean', 'bid_median', 'num_simultaneous_bids', 'num_bids',
'num_auct', 'num_device_type', 'num_url', 'num_ip', 'num_ctry',
'max_bids_per_device', 'num_first_bids', 'num_last_bids'],
dtype='object')

# Calculate various ratios and percentages for train_set
if 'num_simultaneous_bids' not in train_set.columns:
    print("Missing 'num_simultaneous_bids' column in train_set.")

```

```

else:
    train_set['percent_simultaneous_bids'] =
train_set['num_simultaneous_bids'] / train_set['num_bids']
if 'num_bids' not in train_set.columns:
    print("Missing 'num_bids' column in train_set.")
else:
    train_set['bids_per_auct'] = train_set['num_bids'] /
train_set['num_auct']
    train_set['bids_per_device'] = train_set['num_bids'] /
train_set['num_device_type']
    train_set['bids_per_url'] = train_set['num_bids'] /
train_set['num_url']
if 'num_auct' not in train_set.columns:
    print("Missing 'num_auct' column in train_set.")
else:
    train_set['device_per_auct'] = train_set['num_device_type'] /
train_set['num_auct']
if 'num_ip' not in train_set.columns or 'num_ctype' not in
train_set.columns:
    print("Missing 'num_ip' or 'num_ctype' columns in train_set.")
else:
    train_set['ip_per_ctype'] = train_set['num_ip'] /
train_set['num_ctype']
if 'max_bids_per_device' not in train_set.columns:
    print("Missing 'max_bids_per_device' column in train_set.")
else:
    train_set['percent_max_bids'] = train_set['max_bids_per_device'] /
train_set['num_bids']

# Fill any NaN values with 0 for train_set
train_set = train_set.fillna(0)

# Calculate various ratios and percentages for test_set
if 'num_simultaneous_bids' not in test_set.columns:
    print("Missing 'num_simultaneous_bids' column in test_set.")
else:
    test_set['percent_simultaneous_bids'] =
test_set['num_simultaneous_bids'] / test_set['num_bids']

if 'num_bids' not in test_set.columns:
    print("Missing 'num_bids' column in test_set.")
else:
    test_set['bids_per_auct'] = test_set['num_bids'] /
test_set['num_auct']
    test_set['bids_per_device'] = test_set['num_bids'] /
test_set['num_device_type']
    test_set['bids_per_url'] = test_set['num_bids'] /
test_set['num_url']

if 'num_auct' not in test_set.columns:
    print("Missing 'num_auct' column in test_set.")

```

```

else:
    test_set['device_per_auct'] = test_set['num_device_type'] /
test_set['num_auct']

if 'num_ip' not in test_set.columns or 'num_ctry' not in
test_set.columns:
    print("Missing 'num_ip' or 'num_ctry' columns in test_set.")
else:
    test_set['ip_per_ctry'] = test_set['num_ip'] /
test_set['num_ctry']

if 'max_bids_per_device' not in test_set.columns:
    print("Missing 'max_bids_per_device' column in test_set.")
else:
    test_set['percent_max_bids'] = test_set['max_bids_per_device'] /
test_set['num_bids']

# Fill any NaN values with 0 for test_set
test_set = test_set.fillna(0)

import pandas as pd
import matplotlib.pyplot as plt

# Load data
bids = pd.read_csv('bids.csv')

# Identify outliers with only one bid
num_bids = bids.groupby('bidder_id')
['bid_id'].count().reset_index(name='num_bids')
outliers = num_bids[num_bids['num_bids'] == 1]

# Print outliers
print('Outliers:')
print(outliers)

# Remove outliers
bids = bids[~bids['bidder_id'].isin(outliers['bidder_id'])]

# Load train data
train = pd.read_csv('train.csv')
train = train.rename(columns={'outcome': 'bot'})

# Merge bids with train to get outcome for each bidder
bids_with_outcome = bids.merge(train[['bidder_id', 'bot']],
on='bidder_id')

# Aggregate number of bids per bidder with outcome
num_bids_with_outcome = bids_with_outcome.groupby('bidder_id')
['bid_id'].count().reset_index(name='num_bids')


```

```
# Plot histogram
fig, ax = plt.subplots()
num_bids_with_outcome['num_bids'].hist(ax=ax, bins=50)
ax.set_xlabel('Number of bids')
ax.set_ylabel('Number of bidders')
plt.show()
```

Outliers:

	bidder_id	num_bids
0	001068c415025a009fee375a12cff4fcnht8y	1
2	0030a2dd87ad2733e0873062e4f83954mkj86	1
9	009479273c288b1dd096dc3087653499lrx3c	1
16	00dd948c3a88f7b68f1952dbeeac68ffb6qoc	1
27	0176025cc599cb59f825d592b8ef3ee3p5aqv	1
...
6587	fecea7c93f6fc416ab1165267723b0bewb7le	1
6596	ff375d34745157a44e3ba3de0f30dd1bhypuh	1
6600	ff5069626488d0409be146cff3f1f2eak2n7a	1
6611	ffd29eb307a4c54610dd2d3d212bf3bagmmpl	1
6613	fff2c070d8200e0a09150bd81452ce29ngcnv	1

[1057 rows x 2 columns]



```
import pandas as pd
```

```

# Read in the bids and train data
bids = pd.read_csv('bids.csv')
train = pd.read_csv('train.csv')

# Identify the bots and humans
bots = train[train['outcome'] == 1]['bidder_id']
humans = train[train['outcome'] == 0]['bidder_id']

# Filter the bids for bots and humans
bot_bids = bids[bids['bidder_id'].isin(bots)]
human_bids = bids[bids['bidder_id'].isin(humans)]

# Count the occurrences of each merchandise category for bots and humans
bot_merchandise_counts = bot_bids['merchandise'].value_counts()
human_merchandise_counts = human_bids['merchandise'].value_counts()

# Print the merchandise counts for bots and humans
print("Merchandise counts for bots:\n", bot_merchandise_counts)
print("Merchandise counts for humans:\n", human_merchandise_counts)

# Investigate the merchandise feature
print(bids['merchandise'].value_counts())

```

```

Merchandise counts for bots:
sporting goods      230326
mobile              105138
jewelry             37101
home goods          18708
computers           11667
office equipment    7967
books and music     1509
Name: merchandise, dtype: int64
Merchandise counts for humans:
sporting goods      939398
jewelry             555634
mobile              492350
home goods          389249
office equipment    160671
furniture           87807
books and music     13733
auto parts          9757
computers           9733
clothing            476
Name: merchandise, dtype: int64
mobile              2126587
jewelry             1902058
sporting goods      1855207
home goods          1224234
office equipment    289838

```

```
furniture          99181
computers          81084
books and music    51941
clothing           16447
auto parts         9757
Name: merchandise, dtype: int64
```

```
# Load data
```

```
train = pd.read_csv('train.csv')
bids = pd.read_csv('bids.csv')
```

```
# Merge bids with train to get outcome for each bidder
```

```
bids_with_outcome = bids.merge(train[['bidder_id', 'outcome']],
on='bidder_id')
```

```
# Get the top 3 merchandises bidded by humans and bots
```

```
top_merchandises = bids_with_outcome.groupby(['merchandise',
'outcome'])['bid_id'].count().reset_index(name='count')
top_merchandises = top_merchandises.pivot(index='merchandise',
columns='outcome', values='count').fillna(0)
top_merchandises['total'] = top_merchandises.sum(axis=1)
top_merchandises = top_merchandises.sort_values('total',
ascending=False).head(3)
```

```
print(top_merchandises)
```

```
outcome          0.0          1.0          total
merchandise
sporting goods  939398.0  230326.0  1169724.0
mobile         492350.0  105138.0   597488.0
jewelry        555634.0   37101.0   592735.0
```

```
# Compute the number of bids for each bidder
```

```
num_bids_per_bidder = bids.groupby('bidder_id')
['bid_id'].count().reset_index(name='num_bids')
```

```
# Merge with the bids_with_outcome
```

```
bids_with_outcome = bids_with_outcome.merge(num_bids_per_bidder,
on='bidder_id', how='left').fillna(0)
```

```
print(bids.columns)
```

```
Index(['bid_id', 'bidder_id', 'auction', 'merchandise', 'device',
'time',
      'country', 'ip', 'url'],
      dtype='object')
```

```
# Select features to plot
```

```
features = ['num_bids', 'merchandise']
```


```
# Generate density plots for each feature
```


```

# Generate density plot for num_bids
plt.figure()
sns.kdeplot(bids_with_outcome.loc[bids_with_outcome['outcome'] == 0,
'num_bids'], label='Human')
sns.kdeplot(bids_with_outcome.loc[bids_with_outcome['outcome'] == 1,
'num_bids'], label='Bot')
plt.xlabel('num_bids')
plt.legend()
plt.show()

# Generate count plot for merchandise
plt.figure()
sns.countplot(data=bids_with_outcome, x='merchandise', hue='outcome')
plt.xlabel('merchandise')
plt.legend(['Human', 'Bot'])
plt.xticks(rotation=45)
plt.show()

```





Model Validation

```
print(train.columns)
```

```
Index(['bidder_id', 'payment_account', 'address', 'outcome'],
      dtype='object')
```

```
import numpy as np
from sklearn.metrics import roc_auc_score
from sklearn.model_selection import train_test_split
from imblearn.over_sampling import RandomOverSampler
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import GridSearchCV,
RepeatedStratifiedKFold
from sklearn.metrics import classification_report, accuracy_score
from imblearn.pipeline import make_pipeline
```

```
# Create RandomOverSampler object to oversample minority class
oversampler = RandomOverSampler()
```

```
# Create a DataFrame containing only the features
features = train_set[['count', 'bid_mean', 'bid_median',
```

```

'num_simultaneous_bids', 'num_bids',
                        'num_auct', 'num_device_type', 'num_url',
'num_ip', 'num_ctry',
                        'max_bids_per_device', 'num_first_bids',
'num_last_bids',
                        'percent_simultaneous_bids', 'bids_per_auct',
'bids_per_device',
                        'bids_per_url', 'device_per_auct',
'ip_per_ctry', 'percent_max_bids']]

# Replace infinite values with NaN
features = features.replace([np.inf, -np.inf], np.nan)

# Fill missing values with median of respective columns
features_filled = features.fillna(features.median())

# Create a Series containing the target variable ('outcome' column)
target = train_set['outcome']

# Split the data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(features_filled,
target, test_size=0.3, random_state=42)

# Oversample the training data
X_train_resampled, y_train_resampled =
oversampler.fit_resample(X_train, y_train)

# Define a function to perform hyperparameter tuning for a single
Random Forest model
def tune_model(X, y):
    # Define parameter grid for GridSearchCV
    param_grid = {'n_estimators': [50, 100, 150], 'max_depth': [5, 10,
15], 'min_samples_split': [2, 5, 10], 'min_samples_leaf': [1, 2, 4]}
    # Create RandomForestClassifier object
    rf = RandomForestClassifier(random_state=42)
    # Create GridSearchCV object
    grid_search = GridSearchCV(rf, param_grid, cv=5,
scoring='accuracy')
    # Fit GridSearchCV object to data
    grid_search.fit(X, y)
    # Print best hyperparameters and corresponding accuracy score
    print('Best hyperparameters:', grid_search.best_params_)
    print('Accuracy:', grid_search.best_score_)

# Perform hyperparameter tuning for each model
for i in range(5):
    # Set different random_state for each model
    random_state = 42 + i
    # Create RandomForestClassifier object with given random_state

```

```

rf = RandomForestClassifier(random_state=random_state)
# Train the model on the oversampled training data
rf.fit(X_train_resampled, y_train_resampled)
# Make predictions on the testing data
y_pred = rf.predict(X_test)
# Print classification report and accuracy score
print('Model', i+1, 'Classification Report:')
print(classification_report(y_test, y_pred))
print('Model', i+1, 'Accuracy Score:', accuracy_score(y_test,
y_pred))
# Perform hyperparameter tuning for the model
tune_model(X_train_resampled, y_train_resampled)

```

Model 1 Classification Report:

	precision	recall	f1-score	support
0.0	0.97	0.99	0.98	577
1.0	0.60	0.33	0.43	27
accuracy			0.96	604
macro avg	0.78	0.66	0.70	604
weighted avg	0.95	0.96	0.95	604

Model 1 Accuracy Score: 0.9602649006622517

Best hyperparameters: {'max_depth': 15, 'min_samples_leaf': 1, 'min_samples_split': 5, 'n_estimators': 100}

Accuracy: 0.9872490531301163

Model 2 Classification Report:

	precision	recall	f1-score	support
0.0	0.97	0.99	0.98	577
1.0	0.56	0.33	0.42	27
accuracy			0.96	604
macro avg	0.77	0.66	0.70	604
weighted avg	0.95	0.96	0.95	604

Model 2 Accuracy Score: 0.9586092715231788

Best hyperparameters: {'max_depth': 15, 'min_samples_leaf': 1, 'min_samples_split': 5, 'n_estimators': 100}

Accuracy: 0.9872490531301163

Model 3 Classification Report:

	precision	recall	f1-score	support
0.0	0.97	0.99	0.98	577
1.0	0.56	0.33	0.42	27
accuracy			0.96	604
macro avg	0.77	0.66	0.70	604
weighted avg	0.95	0.96	0.95	604

Model 3 Accuracy Score: 0.9586092715231788

Best hyperparameters: {'max_depth': 15, 'min_samples_leaf': 1, 'min_samples_split': 5, 'n_estimators': 100}

Accuracy: 0.9872490531301163

Model 4 Classification Report:

	precision	recall	f1-score	support
0.0	0.97	0.99	0.98	577
1.0	0.54	0.26	0.35	27
accuracy			0.96	604
macro avg	0.75	0.62	0.66	604
weighted avg	0.95	0.96	0.95	604

Model 4 Accuracy Score: 0.956953642384106

Best hyperparameters: {'max_depth': 15, 'min_samples_leaf': 1, 'min_samples_split': 5, 'n_estimators': 100}

Accuracy: 0.9872490531301163

Model 5 Classification Report:

	precision	recall	f1-score	support
0.0	0.97	0.99	0.98	577
1.0	0.54	0.26	0.35	27
accuracy			0.96	604
macro avg	0.75	0.62	0.66	604
weighted avg	0.95	0.96	0.95	604

Model 5 Accuracy Score: 0.956953642384106

Best hyperparameters: {'max_depth': 15, 'min_samples_leaf': 1, 'min_samples_split': 5, 'n_estimators': 100}

Accuracy: 0.9872490531301163

```
import numpy as np
```

```
# Load data
```

```
bids = pd.read_csv('bids.csv')
```

```
# Remove outliers
```

```
outliers =
```

```
bids.groupby('bidder_id').size().reset_index(name='num_bids')
```

```
outliers = outliers[outliers['num_bids'] == 1]['bidder_id']
```

```
bids = bids[~bids['bidder_id'].isin(outliers)]
```

```
# Use train_set DataFrame
```

```
train = train_set
```

```
# Replace infinite values with NaN
```

```

train = train.replace([np.inf, -np.inf], np.nan)

# Fill missing values with the mean
train = train.fillna(train.mean())

# Define features
features = ['count', 'bid_mean', 'bid_median',
            'num_simultaneous_bids', 'num_bids',
            'num_auct', 'num_device_type', 'num_url', 'num_ip',
            'num_ctry',
            'max_bids_per_device', 'num_first_bids', 'num_last_bids',
            'percent_simultaneous_bids', 'bids_per_auct',
            'bids_per_device',
            'bids_per_url', 'device_per_auct', 'ip_per_ctry',
            'percent_max_bids']

# Separate features and target
X = train[features]
y = train['outcome']

# Perform over-sampling
ros = RandomOverSampler(random_state=42)
X_resampled, y_resampled = ros.fit_resample(X, y)

# Initialize base models
rf1 = RandomForestClassifier(random_state=1)
rf2 = RandomForestClassifier(random_state=2)
rf3 = RandomForestClassifier(random_state=3)

# Calculate AUC for each base model
for model in [rf1, rf2, rf3]:
    model.fit(X_resampled, y_resampled)
    y_pred = model.predict_proba(X)[ :, 1]
    auc = roc_auc_score(y, y_pred)
    print(f"AUC for model with random state {model.random_state}:
{auc}")

```

C:\Users\egodd\AppData\Local\Temp\ipykernel_30568\797303328.py:18:
FutureWarning: The default value of numeric_only in DataFrame.mean is
deprecated. In a future version, it will default to False. In
addition, specifying 'numeric_only=None' is deprecated. Select only
valid columns or specify the value of numeric_only to silence this
warning.

```
train = train.fillna(train.mean())
```

```

AUC for model with random state 1: 1.0
AUC for model with random state 2: 1.0
AUC for model with random state 3: 1.0

```

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import roc_auc_score
```

```
# Initialize the model
```

```
rfc = RandomForestClassifier(random_state=42)
```

```
# Fit the model on the training data
```

```
rfc.fit(X_train_resampled, y_train_resampled)
```

```
# Predict the probabilities of class 1 (bot)
```

```
y_pred_prob = rfc.predict_proba(X_test)[:, 1]
```

```
# Compute the AUC score
```

```
auc_score = roc_auc_score(y_test, y_pred_prob)
```

```
print("AUC score:", auc_score)
```

```
AUC score: 0.8899159124462417
```

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import roc_auc_score
```

```
# Initialize the models with different random states
```

```
rfc1 = RandomForestClassifier(random_state=42)
```

```
rfc2 = RandomForestClassifier(random_state=10)
```

```
rfc3 = RandomForestClassifier(random_state=123)
```

```
# Fit the models on the training data
```

```
rfc1.fit(X_train_resampled, y_train_resampled)
```

```
rfc2.fit(X_train_resampled, y_train_resampled)
```

```
rfc3.fit(X_train_resampled, y_train_resampled)
```

```
# Predict the probabilities of class 1 (bot) for each model
```

```
y_pred_prob1 = rfc1.predict_proba(X_test)[:, 1]
```

```
y_pred_prob2 = rfc2.predict_proba(X_test)[:, 1]
```

```
y_pred_prob3 = rfc3.predict_proba(X_test)[:, 1]
```

```
# Ensemble averaging by taking the mean of the probabilities
```

```
y_pred_prob = (y_pred_prob1 + y_pred_prob2 + y_pred_prob3) / 3
```

```
# Compute the AUC score
```

```
auc_score = roc_auc_score(y_test, y_pred_prob)
```

```
print("AUC score:", auc_score)
```

```
AUC score: 0.8878297708453687
```

```
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import roc_auc_score
from sklearn.model_selection import GridSearchCV
```

```
# Initialize the model
```

```

rfc = RandomForestClassifier(random_state=42)

# Define the hyperparameter grid
param_grid = {
    'n_estimators': [50, 100, 200],
    'max_depth': [5, 10, 15],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 5]
}

# Initialize GridSearchCV with 5-fold cross-validation
grid_search = GridSearchCV(rfc, param_grid, cv=5, scoring='roc_auc')

# Fit GridSearchCV on the training data
grid_search.fit(X_train_resampled, y_train_resampled)

# Print the best hyperparameters and AUC score
print("Best hyperparameters:", grid_search.best_params_)
print("AUC score:", grid_search.best_score_)

Best hyperparameters: {'max_depth': 15, 'min_samples_leaf': 1,
'min_samples_split': 2, 'n_estimators': 50}
AUC score: 1.0

import numpy as np
from sklearn.model_selection import StratifiedKFold

# Initialize StratifiedKFold with 5 splits
skf = StratifiedKFold(n_splits=5, random_state=42, shuffle=True)

# Create empty lists to store the ensemble_scores and
grid_search_scores
ensemble_scores = []
grid_search_scores = []

# Loop through the splits created by StratifiedKFold
for train_index, test_index in skf.split(X, y):
    X_train, X_test = X.iloc[train_index], X.iloc[test_index]
    y_train, y_test = y.iloc[train_index], y.iloc[test_index]

    # Perform over-sampling
    X_train_resampled, y_train_resampled = ros.fit_resample(X_train,
y_train)

    # Ensemble averaging
    rfc1.fit(X_train_resampled, y_train_resampled)
    rfc2.fit(X_train_resampled, y_train_resampled)
    rfc3.fit(X_train_resampled, y_train_resampled)

    y_pred_prob1 = rfc1.predict_proba(X_test)[: , 1]

```

```

y_pred_prob2 = rfc2.predict_proba(X_test)[: , 1]
y_pred_prob3 = rfc3.predict_proba(X_test)[: , 1]

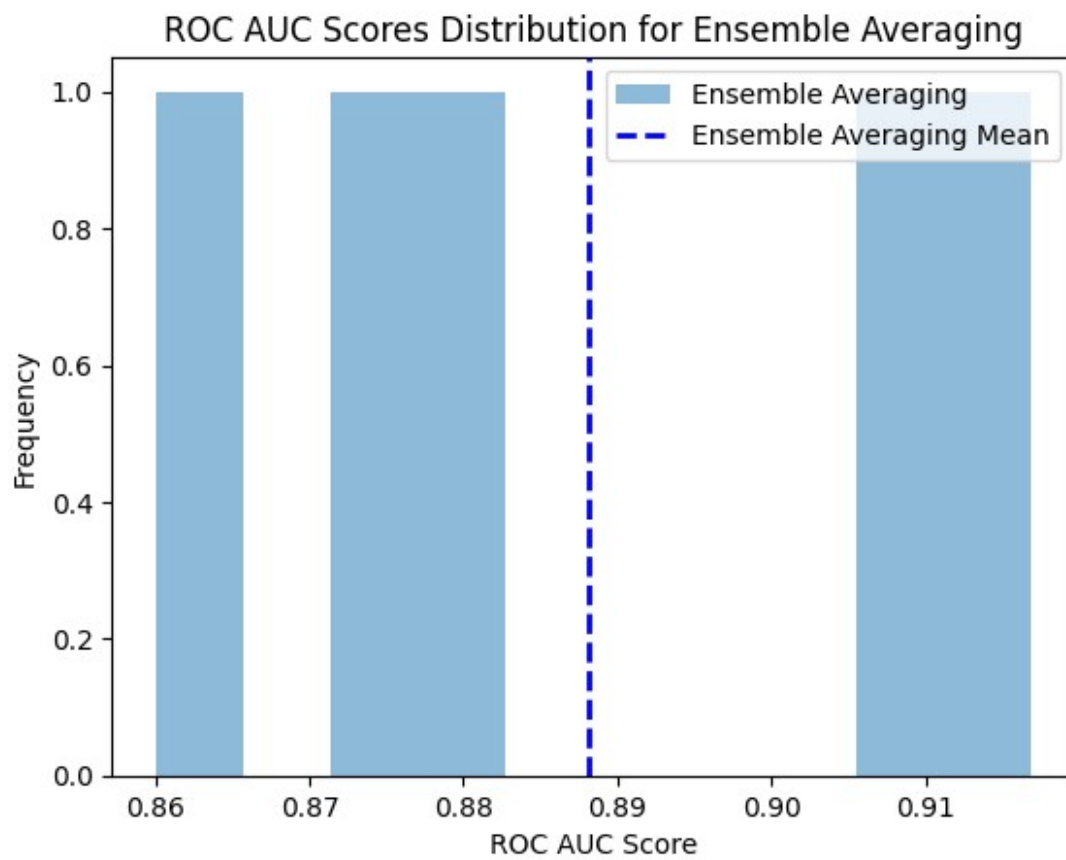
y_pred_prob = (y_pred_prob1 + y_pred_prob2 + y_pred_prob3) / 3
ensemble_auc_score = roc_auc_score(y_test, y_pred_prob)
ensemble_scores.append(ensemble_auc_score)


# Hyperparameter tuning
grid_search.fit(X_train_resampled, y_train_resampled)
best_model = grid_search.best_estimator_
y_pred_prob = best_model.predict_proba(X_test)[: , 1]
grid_search_auc_score = roc_auc_score(y_test, y_pred_prob)
grid_search_scores.append(grid_search_auc_score)

# Plot ROC AUC score distribution for ensemble averaging
plt.hist(ensemble_scores, alpha=0.5, label='Ensemble Averaging')
plt.axvline(x=np.mean(ensemble_scores), color='b', linestyle='dashed',
linewidth=2, label='Ensemble Averaging Mean')
plt.title('ROC AUC Scores Distribution for Ensemble Averaging')
plt.xlabel('ROC AUC Score')
plt.ylabel('Frequency')
plt.legend()
plt.show()

# Plot ROC AUC score distribution for hyperparameter tuning
plt.hist(grid_search_scores, alpha=0.5, label='Hyperparameter Tuning')
plt.axvline(x=np.mean(grid_search_scores), color='r',
linestyle='dashed', linewidth=2, label='Hyperparameter Tuning Mean')
plt.title('ROC AUC Scores Distribution for Hyperparameter Tuning')
plt.xlabel('ROC AUC Score')
plt.ylabel('Frequency')
plt.legend()
plt.show()

```



Final Model

Define the final models

```
final_models = [RandomForestClassifier(n_estimators=100, max_depth=8,
                                       random_state=42),
                RandomForestClassifier(n_estimators=200, max_depth=10,
                                       random_state=123),
                RandomForestClassifier(n_estimators=300, max_depth=12,
                                       random_state=456)]
```

Fit the models and store their predictions on the test set

```
proba = []
for model in final_models:
    model.fit(X_resampled, y_resampled)
    proba_rforest = model.predict_proba(X_test)[:,-1]
    proba.append(proba_rforest)
```

Average the predictions of the models

```
result = np.mean(proba, axis = 0)
```

Create the output dataframe and save it to a CSV file

```
output_dataframe = pd.DataFrame({
    'bidder_id': test_set['bidder_id'].iloc[:len(result)],
    'prediction': result
})
```

```
})  
output_dataframe.to_csv('my_predictions.csv', index=False)
```